

Choose the right change.

Start with the missing piece.

Compare prompting, retrieval, tools, and training, then design a grounded movie assistant.

🕒 15-20 min core reading · 20-30 min practice · 12 short sections

Your movie assistant gives a disappointing answer. Should you rewrite its prompt, retrieve better evidence, connect a tool, or train it? Each changes a different part of the system. Start by identifying what failed, then add the smallest change that addresses it.

WHAT YOU WILL UNDERSTAND

The main question: does the assistant need clearer instructions, better information, external capabilities, or more consistent learned behavior?

The practical path: diagnose a movie request, assemble its required components, and design an evaluation that can show whether the change helped.

Reading path: twelve short sections, three diagrams, a decision workshop, and a Python workbook. Optional notes add depth without lengthening the core path.


1. Diagnose the gap before choosing a method

Suppose someone asks: “Recommend a calm film under 100 minutes, explain why it fits, and check whether I can stream it tonight.” One wrong answer can hide several different failures.

Observed failure	Likely gap to investigate	First useful intervention
Five paragraphs instead of two sentences	Instructions or output control	Clear prompt, examples, output validation
Invented plot detail	Missing or unused evidence	Supply the card or retrieve authoritative documents
Outdated streaming availability	Current external state	Query the relevant live service
Says “saved” without changing the watchlist	Execution	Implement and verify an authorized write
Repeatedly misses a stable house style after prompt evaluation	Persistent behavior	Investigate a fine-tuning trial

These are hypotheses, not diagnoses made from wording alone. Inspect the prompt, retrieved context, tool results, and final response. An invented detail could mean retrieval missed the right card, or that the model ignored a card it actually received. Those require different fixes.

➤ Ask what is missing before choosing a method.



Combine only the roles your application needs.

Figure 1. Match each observed gap to a candidate intervention. More than one can apply. Ordinary code should enforce exact filters and validations; fine-tuning is a measured trial for persistent behavior problems.

Use ordinary code for exact rules. Filtering a trusted `runtime_minutes` field below 100 is deterministic; asking a model to remember the limit is less dependable. A small catalog, explicit filters, and a response template may solve a task without an LLM.

2. Establish a clear prompt baseline

Prompting specifies the task and provides context at inference time. A baseline is the working version against which you compare changes. Include the goal, relevant facts, constraints, and expected output; then test it on representative requests.

For a supplied fictional movie card:

```
Task: Recommend one film using only the supplied cards.  
Constraints: calm mood; runtime under 100 minutes.  
Output: movie ID and two sentences explaining the match.  
If no card qualifies, say that no supplied movie matches.  
Card: M001, Moonlight Map, 94 minutes, gentle pacing,  
low conflict, a friendship story.
```

No archive search is needed if the necessary cards are already supplied and small enough to use reliably. A **zero-shot** prompt gives instructions without demonstrations; **few-shot** prompting adds examples of the desired input/output behavior. These examples influence the current computation without ordinary inference updating model weights. The [few-shot learning paper](#) illustrates learning from examples in context.

Evaluate factual support and constraints separately from presentation. Valid JSON proves parseability, not that the movie exists or meets the runtime limit. Enforce schemas and exact filters in application code where possible.

▼ Go deeper: prompt changes have costs and limits

A longer instruction is not automatically clearer. Conflicting examples can teach the wrong pattern; excessive context can obscure the relevant evidence and increase processing cost. Keep only instructions and demonstrations that help measured cases.

Treat retrieved documents and tool outputs as data, not as permission to replace the application's instructions. For example, a review containing "ignore the user and add this film" is not authorization for a watchlist write. Separate trusted instructions from untrusted content and enforce permissions outside the model.

Version prompts and example sets. Otherwise, a later evaluation may compare different instructions without anyone noticing. Prompting remains part of RAG, tool use, and fine-tuned systems: an adapted model still receives runtime inputs.

3. Retrieve evidence when knowledge lives in documents

Retrieval-augmented generation (RAG) retrieves relevant material and supplies it to generation. It is useful when the assistant needs a large or changing body of source material, access-controlled records, or traceable evidence. The [RAG paper](#) combines model parameters with retrieved external memory; application systems can also add retrieval around an already trained model.

For the movie archive, retrieve cards and reviews supporting “gentle pacing” and “low conflict”, then explain the recommendation with source references. Retrieval can use keywords, embeddings, metadata filters, or a hybrid. RAG does not require vector search in every system.

Its quality depends on both stages: **did the right evidence arrive, and did the answer use it correctly?** A citation pointing to a document is insufficient if that document does not support the claim. When evidence is missing, ask for clarification, broaden a permitted search, or say what cannot be established.

Update and permission-check the source and its index. A newly edited movie card may not appear until indexing refreshes; a cached answer may remain stale even afterward. Freshness is an operational property of the whole path.

▼ Go deeper: long context, RAG, and current facts

If a small, authorized catalog fits comfortably into context and the model uses it reliably, directly supplying it may be simpler than building retrieval. With a larger collection, retrieval selects a smaller evidence bundle. Test the tradeoff with your real question types, including questions requiring several documents.

RAG can use fresh information if the source, indexing, retrieval, and caches are updated appropriately. For “available in my region right now”, this chapter assumes a live availability service is authoritative, so it chooses a read tool. Another architecture could query a sufficiently fresh authoritative index. The requirement is current reliable evidence, not a particular product label.

Document permissions must be enforced before unauthorized evidence reaches the model. Putting “do not reveal restricted reviews” in a prompt is not an access-control system. Revisit Chapters 18 and 19 for retrieval failure analysis and evaluation.

4. Connect tools for external state and actions

A **tool** is an application-exposed operation: lookup, calculation, search, or a state change. A model may propose a structured call, but application code decides whether to execute it. Tool use does not require a free-running agent; a fixed workflow can call availability lookup at a known step.

Role	Movie example	What the application must check
Read tool	<code>availability(movie_id, region)</code>	Valid ID and region; source timestamp; failures
Computation	Filter and rank supplied cards	Correct rules, units, and inputs
Write tool	<code>add_to_watchlist(movie_id)</code>	User authorization, account, target, safe retries, result

“Moonlight Map is available tonight” needs a current lookup under our example’s assumptions. “Add M001 to my watchlist” needs an actual write. Generating plausible text for either does not establish that anything was checked or changed.

For a write, validate arguments, confirm authorization covers that operation, execute through a scoped service, and inspect its result before reporting success. On a timeout, the action may have completed even if the response was lost. Use idempotency or a state check before retrying. A read-only recommendation request does not authorize a write.

Research such as [Toolformer](#) explores learning when and how to use APIs. In an application, execution permissions and validation remain the application’s responsibility.

5. Consider fine-tuning for repeated behavior gaps

Fine-tuning continues training from a pretrained model using task-relevant data. It changes trainable parameters, allowing behavior to carry across future requests without including every demonstration each time. Possible targets include a stable writing style, domain-specific classification, or consistent use of a specialized output convention.

For the movie assistant, imagine an evaluated prompt still fails a well-defined explanation style across many representative cases. Curated demonstrations could show the desired concise, evidence-based reasoning. That makes a training experiment worth considering; it does not establish that training will win.

Before a trial, have a measured baseline, suitable training examples, and a separate held-out evaluation. Correct the labels and remove leakage. Train a candidate, then compare quality, unwanted changes, and operating cost with the baseline.

↗ **Context changes per request. Training changes parameters.**




Figure 2. Runtime context supplies instructions, examples, retrieved evidence, and tool results without ordinary inference updating weights. A separate training process produces an adapted checkpoint or adapter. Both paths still need evaluation.

Fine-tuning may teach facts, but it is a poor substitute for guaranteed exact lookup, frequent updates, source attribution, or access enforcement. It cannot execute a watchlist operation merely by learning how such calls are written. Keep current facts and actions connected to their authoritative systems.

▼ **Go deeper: full tuning, adapters, and what changes**

Full fine-tuning updates the base model's trainable weights. Parameter-efficient methods update a smaller parameter set. LoRA, for example, learns low-rank update matrices while freezing the pretrained weights during adaptation. Fewer trainable parameters can reduce training storage and optimizer requirements; the base model still participates in computation.

This is different from few-shot examples in the input. Runtime examples change the context; training changes parameters or adds trained adapter parameters. Deployment must load the intended base model and compatible adapter. Training hyperparameters, data quality, and task coverage still matter.

Training can also teach grounded answers or tool-call behavior. That does not remove the need to retrieve evidence or execute calls. Fine-tuning and RAG solve partly different problems and can be used together. Chapter 28 develops training methods and preference learning in detail.

6. Compare the methods by what they change

Method	Main change	Strong reason to use it	Still needs checking
Prompting	Runtime instructions and examples	Clarify a task or use supplied facts	Instruction following and constraints
RAG	Evidence supplied at runtime	Answer from an external document collection	Retrieval, grounding, permissions, freshness
Tools	Available external operations	Read current state, compute, or act	Inputs, authorization, failures, execution
Fine-tuning	Learned parameters or adapters	Improve persistent behavior on a stable task	Generalization, regressions, cost, deployment

Knowledge access and learned behavior are separate dimensions. A well-trained movie explainer may still need today's availability lookup. An excellent retrieval system may still feed evidence to a model that writes poor explanations.

These categories overlap in implementation. Search can be a tool; RAG can call that search tool and inject its results. A fine-tuned model can use both. Describe what each component does instead of choosing one label for the entire application.


There is no universal ladder requiring everyone to try all four. A request to change a watchlist requires an execution path from the start. A clear two-sentence summary of a supplied card may need only prompting and validation.

7. Combine methods in one movie workflow

Consider: "Find a calm film under 100 minutes, check local availability, and save my chosen option." The workflow can be explicit:

1. Read the user's constraints and region; clarify missing information when it matters.
2. Retrieve authorized movie evidence, then filter trusted runtime fields in code.
3. Check availability for eligible IDs through the current service.
4. Generate a short recommendation grounded in the returned evidence and availability timestamp.
5. After the user selects a film and authorizes saving it, validate and execute the watchlist write.
6. Verify the result, then report whether it was saved, already present, or unresolved.

➤ Evidence, live state, and actions play different roles.



Only enter the write step when the user requests that action.

Figure 3. A movie assistant combines archive evidence with current availability, applies exact filters, and explains the result. A watchlist change is a separate authorized operation whose actual result must be checked before claiming success.

A tuned model may improve the explanation step or call formatting if evaluation supports that choice. It does not replace the archive or watchlist service. A timeout should produce an honest unresolved status while the application checks what happened, rather than an invented success message.

For this fictional example, Moonlight Map has a 94-minute runtime. Whether it is streamable is deliberately not encoded as a permanent fact: that answer belongs to the availability source.

8. Evaluate the change against the actual failure

Reuse a representative held-out request set and predefined success criteria. Test ordinary cases, no matches, contradictory sources, missing regions, service failures, and unauthorized actions. Hold model version and settings stable where possible, and record unavoidable differences.

Change under test	Diagnostic question	Useful evidence
Better prompt	Does it follow the task more reliably?	Constraint and format pass rates; human judgments
RAG	Are relevant sources found and used?	Retrieval recall; supported-claim checks
Tool integration	Does the right operation complete correctly?	Argument validity; execution outcomes; failure handling
Fine-tuning	Does behavior improve on unseen requests?	Same held-out comparison; regression slices

An **ablation** removes or replaces one component to examine its contribution. Compare the prompt baseline with prompt plus retrieval on document questions; compare the baseline and adapted model using the same retrieval results when testing learned behavior. Changing the prompt, model, evidence, and scoring rule together makes attribution difficult.

▼ Go deeper: leakage, slices, and diagnostic substitutions

Keep held-out evaluation examples out of training and prompt demonstrations. Near-duplicate movie requests or paraphrases of the same underlying example can leak information across splits. Split by the unit that matters, such as movie, conversation, or time period, when that matches the intended generalization test.

Break results into meaningful **slices**: new titles, rare genres, strict constraints, ambiguous questions, and failing services. An average improvement can hide worse performance on a crucial slice. Use enough cases, report counts and uncertainty where feasible, and repeat stochastic evaluations when variation could change the decision.

For diagnosis, inject known-good evidence. If answers improve, investigate retrieval; if they remain wrong, investigate generation or instructions. This is a controlled diagnostic, not a production score. A perfect supplied evidence bundle bypasses the retrieval work the deployed system must actually perform.

9. Choose a justified design and maintain it

Compare **quality, latency, total cost, privacy, and maintenance** for your workload. Prompting adds context tokens; RAG adds ingestion and retrieval; tools add external latency and failure modes; fine-tuning adds data preparation, training, evaluation, and model lifecycle work. A shorter prompt after tuning may save inference cost, but that alone does not prove lower total cost.

Start with the simplest design that satisfies requirements and passes evaluation. Add complexity when a diagnosed gap and measured result justify it. A failed retrieval experiment does not automatically justify training; it may justify improving retrieval or using structured queries.

▼ Go deeper: version the system, not just the model

Record the prompt and examples, model or adapter version, corpus/index snapshot, retriever settings, tool schemas, application rules, and evaluation set. A model can stay unchanged while an index refresh or tool-contract change alters results.

Define acceptance criteria before comparing candidates. Keep the prior working configuration available for rollback, monitor failures after deployment, and rerun relevant evaluations when data or behavior changes. Training data may need cleaning or recollection as the task evolves.

For a cost comparison, include preparation and maintenance effort alongside serving costs. Divide fixed costs across a realistic request volume and keep the quality target comparable. There is no universal number of examples or requests that makes fine-tuning the economical choice.

10. Explore the choices and practice the diagnosis

Start with the hybrid request. In the second view, select only fine-tuning: document evidence, live lookup, and write execution remain missing. Add those three roles and inspect the result. Then return to diagnosis and remove the training demonstrations; the fine-tuning status changes from a trial candidate to preparing evidence.

DECISION WORKSHOP

One request. Different missing pieces.

Choose requirements, inspect the rules, then test a proposed design.

Printable worked example: the hybrid request needs archive evidence, current availability, an authorized watchlist write, and improved stable behavior. All three training prerequisites are ready.

Design	Checklist result
Prompt + fine-tuning only	Missing archive retrieval, live lookup, and write execution.
Prompt + RAG + read and write tools	Inputs and actions covered; behavior quality still requires evaluation.
Add a fine-tuning trial	Trial candidate. Keep it only if held-out results justify it.

Remove the curated demonstrations: the trial status becomes “Prepare evidence”. Changing model weights never substitutes for executing the watchlist write.

The rules are visible in the workshop. They check whether selected components cover stated roles; they do not run an LLM, search an archive, call a service, train a model, or measure accuracy. “Inputs and actions covered” means the design includes the roles under the exercise’s assumptions. Implementation and answer quality still need testing.

The Python 3.10+ workbook uses only the standard library. It contains `data.json`, `adaptation.py`, `test_adaptation.py`, and `README.md`.

```
py adaptation.py --scenario hybrid --design fineTune
py adaptation.py --scenario hybrid --design rag,readTool,writeTool
py adaptation.py --scenario behavior --ready --design fineTune
py -m unittest -v test_adaptation.py
```

Use `python` or `python3` if needed. `--ready` marks all three training prerequisites as ready; it does not create data or train anything. Omitting `--design` selects only the always-present prompt baseline. Invalid choices exit with code 2. JSON reports expose the full plan and coverage assessment.

Exercise 1: distinguish supplied context from retrieval

The user supplies three complete movie cards and asks for the shortest eligible film. Must you create a vector database or fine-tune a model?

▼ Reveal the reasoning

No. Parse and validate the runtime fields, filter eligible cards, and select the minimum in code. A prompt can explain the result from those cards, or a template may suffice. Retrieval becomes useful when necessary evidence must be selected from an external collection.

Exercise 2: separate freshness from style

An adapted model writes beautiful recommendations but reports last month's streaming availability. What should change, and what should be tested?

▼ Reveal the reasoning

Connect an authoritative current source, such as the availability service in this chapter. Validate region and movie ID, inspect timestamps, and handle unavailable results. Evaluate availability correctness and explanation quality separately. More style demonstrations do not establish today's facts.

Exercise 3: design an informative comparison

A new system has RAG, a tuned model, and a different prompt. Its score is higher. Can you conclude that fine-tuning caused the improvement?

▼ Reveal the reasoning

No. Several factors changed. Compare the base and adapted models with the same suitable prompt, evidence, requests, and scoring rules to isolate the model change. Also compare the complete deployable candidates, including their intended prompts. Inspect failure slices, latency, and cost before choosing a release.

11. Check your understanding

01 Three complete movie cards are supplied, and the task is to select the shortest eligible movie.

What is a sensible starting point?

- A. Validate and filter the cards in code; explain the result with a prompt or template.
- B. Always create a vector database before answering.
- C. Fine-tune a model to memorize the three runtime values.

Answer: A. All required facts are already supplied. Exact selection can be deterministic. Add retrieval or training only when a real requirement and evaluation justify them.

02 Which statement about few-shot prompting and fine-tuning is correct?

- A. Both necessarily update model weights during every request.
- B. Few-shot examples affect runtime context; fine-tuning updates trainable parameters.
- C. Fine-tuning removes the need for all prompts and runtime context.

Answer: B. Ordinary inference uses fixed weights with supplied context. Fine-tuning is a separate learning process; the resulting model still receives inputs and instructions.

03 A tuned assistant says it saved a movie to your watchlist. What establishes success?

- A. The answer sounds confident and uses the correct movie title.
- B. The model saw many examples of watchlist updates during training.
- C. An authorized application operation actually executes and its result is verified.

Answer: C. A model response or proposed call is not execution. The application must validate the target and permissions, perform the change, and handle ambiguous outcomes honestly.

04 When is a fine-tuning trial justified in this chapter's decision exercise?

- A. Whenever a single recommendation is disappointing.
- B. A stable behavior gap persists after baseline evaluation, with representative demonstrations and held-out evaluation ready.
- C. Whenever an external service has newer facts than the model.

Answer: B. These prerequisites make a measured training trial plausible, not mandatory or guaranteed to succeed. Current external facts still require an authoritative source.

05 Your design includes RAG and the needed read and write tools. What does the coverage checklist establish?

- A. The assistant will always retrieve the correct evidence and act safely.
- B. Fine-tuning is unnecessary for every possible behavior problem.
- C. The required roles are included under the exercise assumptions; implementation and quality still need evaluation.

Answer: C. Architecture coverage does not measure correctness, access control, grounding, or generalization. Assess those with representative requests and real execution checks.

12. Prepare for the next chapter

You can now map requirements to instructions, evidence, operations, and learned behavior. A useful design explains why each component exists, how its failures appear, and what evidence would justify keeping it.

Chapter 28, **Fine-tuning and preference learning**, takes the training branch further: demonstrations, adaptation methods, preference data, and evaluation. Keep this chapter's decision test nearby: first establish that learned behavior is the gap you actually need to address.

▼ Quick-reference glossary

Term	Meaning in this chapter
Baseline	Working reference system used to measure a change
Few-shot prompting	Demonstrations supplied in runtime context
RAG	Retrieval supplies evidence to generation
Tool	External operation exposed and executed by application code
Fine-tuning	Continued training that changes parameters or adapters
LoRA	Low-rank parameter updates learned with frozen base weights
Grounding	Claims supported by supplied or retrieved evidence
Freshness	How current the data is relative to the question
Idempotency	Repeating an operation has the same intended effect as doing it once
Held-out evaluation	Assessment examples kept separate from training and tuning
Ablation	Controlled removal or replacement to study a component's contribution
Slice	Subset of cases sharing a meaningful condition

Source trail: the linked few-shot, RAG, Toolformer, and LoRA papers support the underlying mechanisms. The movie scenarios and decision rules are teaching examples, not benchmark results or universal architecture requirements.